

# Lab. – Classification

## Convolutional Neural Networks (CNNs)

Yuan-Fu Liao

National Yang Ming Chiao Tung University

[yfliao@nycu.edu.tw](mailto:yfliao@nycu.edu.tw)

# Kaggle Competition - The Simpsons Characters Recognition Challenge

<https://www.kaggle.com/t/2d96605587334fdc9eb0fabf6756392f>

- Images of 20 characters extracted from The Simpsons episodes
  - ✓ About 1000 images per character
  - ✓ Pictures are under various size, scenes
  - ✓ not necessarily centered in each image and could sometimes be with or cropped from other characters





InClass Prediction Competition

# Machine Learning@NTUT - Classification

## The Simpsons Characters Recognition Challenge

18 days to go

Overview

**Data**

Discussion

Leaderboard

Rules

Team

Host

My Submissions

**Submit Predictions**

### Competition Data

[Edit](#)

sampleSubmission.csv

20\_characters\_illust...

character\_name.lst

test.zip


**train.zip****train.zip** 482.24 MB **Download**

URL for Sharing :

<https://www.kaggle.com/t/2d96605587334fdc9eb0fabf6756392f>

# Machine Learning@NTUT - 2017

MachineLearningNTUT



## Classification

Individual assignment

**Give this to your students**

<https://classroom.github.com/a/4JnaHLk8>



# 20 Characters in The Simpsons



abraham\_grampa\_simpson  
comic\_book\_guy

apu\_nahasapeemapetilon  
edna\_krabappel

homer\_simpson

bart\_simpson  
kent\_brockman

charles\_montgomery\_burns  
krusty\_the\_clown

chief\_wiggum

lenny\_leonard  
moe\_szyslakned\_flanders

lisa\_simpson

nelson\_muntz

marge\_simpson

principal\_skinner

mayor\_quimby

sideshow\_bob  
milhouse\_van\_houten



# Demo






- Task1
  - Predict Simpsons Characters in all pictures



- Task2

- **Compute the Confusion Matrix**






# • Task3

## • Visualization and Understanding Convolutional Neural Networks

- 畫出每一層filter的權重 只需要抓第一層的權重，tensor轉image
- ~~畫出每一層的feature map~~



# Example

- Tensorflow Tutorial - Convolutional Neural Networks
  - [https://www.tensorflow.org/versions/r0.11/tutorials/deep\\_cnn/index.html](https://www.tensorflow.org/versions/r0.11/tutorials/deep_cnn/index.html)
  - The code for this tutorial resides in [tensorflow/models/image/cifar10/](#)
  - Accuracy on Test-Set: 79.3% (7932 / 10000)

# CIFAR-10

- CIFAR-10 classification is a common benchmark problem in machine learning.
- The problem is to classify RGB 32x32 pixel images across 10 categories: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck

**airplane**



**automobile**



**bird**



**cat**



**deer**



**dog**



**frog**



**horse**



**ship**




**truck**





# AlexNet





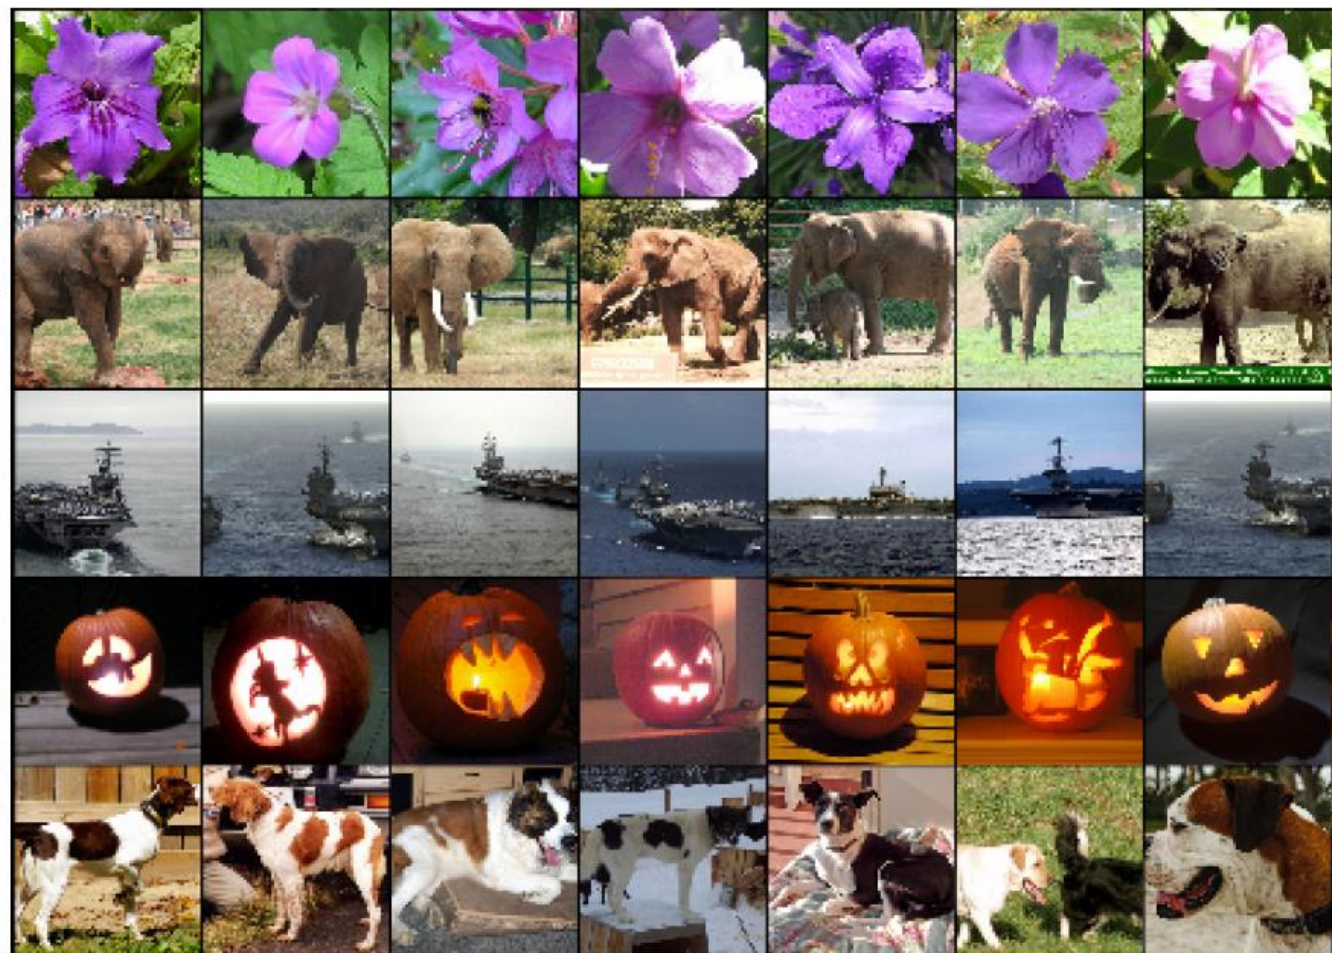
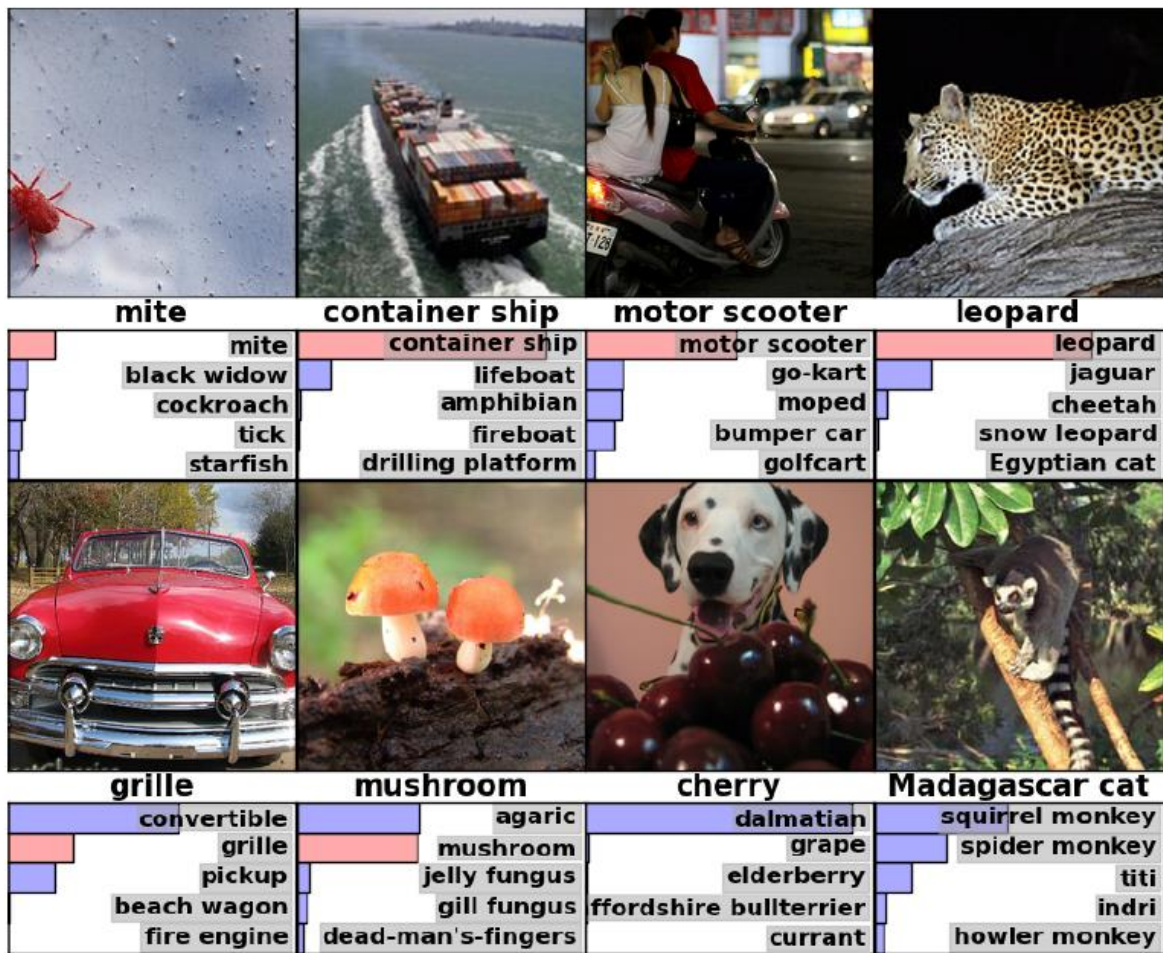


Figure 4: **(Left)** Eight ILSVRC-2010 test images and the five labels considered most probable by our model. The correct label is written under each image, and the probability assigned to the correct label is also shown with a red bar (if it happens to be in the top 5). **(Right)** Five ILSVRC-2010 test images in the first column. The remaining columns show the six training images that produce feature vectors in the last hidden layer with the smallest Euclidean distance from the feature vector for the test image.



# References

- TensorFlow Tutorial CIFAR-10
  - [https://www.youtube.com/watch?v=3BXfw\\_1\\_TF4](https://www.youtube.com/watch?v=3BXfw_1_TF4)
- TensorFlow CIFAR-10 tutorial, detailed step-by-step review
  - Part 1: <http://www.aimechanic.com/2016/10/13/d242-tensorflow-cifar-10-tutorial-detailed-step-by-step-review-part-1/>
  - Part 2: <http://www.aimechanic.com/2016/10/17/d246-tensorflow-cifar-10-tutorial-detailed-review-part-2/>
- [TensorFlow-Tutorials/06\\_CIFAR-10.ipynb](https://github.com/Hvass-Labs/TensorFlow-Tutorials/blob/master/06_CIFAR-10.ipynb)
  - [https://github.com/Hvass-Labs/TensorFlow-Tutorials/blob/master/06\\_CIFAR-10.ipynb](https://github.com/Hvass-Labs/TensorFlow-Tutorials/blob/master/06_CIFAR-10.ipynb)




# ImageNet Classification with Deep Convolutional Neural Networks

# Goa

I



Classifica(on



# ImageNet

- Over 15M labeled high resolution images
- Roughly 22K categories
- Collected from web and labeled by Amazon Mechanical Turk





# ILSVR C

- Annual competition of image classification at large scale
- 1.2M images in 1K categories
- Classification: make 5 guesses about the image label




EntleBucher



Appenzeller


# Convolutional Neural Networks

- Model with a large learning capacity
- Prior knowledge to compensate all data we do not have




# ILSVRC

ImageNet Classification error throughout years and groups




Li Fei-Fei: ImageNet Large Scale Visual Recognition Challenge, 2014 <http://image-net.org/>

# ILSVRC Top 5 Error on ImageNet





# ImageNet Classification Error (Top 5)




# SuperVision (SV)


Image classification with deep convolutional neural networks

- 7 hidden “weight” layers
  - 650K neurons
  - 60M parameters
  - 630M connections
- 
- Rectified Linear Units, overlapping pooling, dropout trick
  - Randomly extracted 224x224 patches for more data

# Architecture




# Layer 1 (Convolutional)



- Images: 227x227x3
- F (receptive field size): 11
- S (stride) = 4
- Conv layer output:  
55x55x96



# Layer 1 (Convolutional)



- $55 \times 55 \times 96 = 290,400$  neurons
- each has  $11 \times 11 \times 3 = 363$  weights and 1 bias
- $290400 \times 364 = 105,705,600$  parameters on the first layer of the AlexNet alone!


# Architecture

## ReLU Nonlinearity

- Standard way to model a

neuron  $f(x) = \tanh(x)$  or  $f(x) = (1 + e^{-x})^{-1}$


Very slow to  
train



- Non-saturating nonlinearity (ReLU)


$$f(x) = \max(0, x)$$

Quick to train



# Architecture

ReLU  
Nonlinearity



A 4 layer CNN with ReLUs (solid line) converges **six times faster** than an equivalent network with tanh neurons (dashed line) on CIFAR-10 dataset




# Architecture

Training on Multiple

GPUs

GPU #1

intra-GPU connec(ons




GPU #2

inter-GPU connec(ons

# Architecture

## Training on Multiple GPUs

GPU #1



GPU #2


inter-GPU connections

Top-1 and Top-5 error rates decreases by 1.7% & 1.2% respectively, comparing to the net trained with one GPU and half neurons!!

# Architecture

Overlapping  
Pooling

Max-pooling layers



Response normalization layers



# Architecture

## Local Response Normalization

- No need to input normalization with ReLUs.
- But still the following local normalization scheme helps generalization.

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Response-normalized activity


Activity of a neuron computed by applying kernel  $l$  at position  $(x,y)$  and then applying the ReLU nonlinearity

- Response normalization reduces top-1 and top-5 error rates by 1.4% and 1.2% , respectively.

# Architectur


## Overlapping Pooling

- Traditional pooling ( $s = z$ )

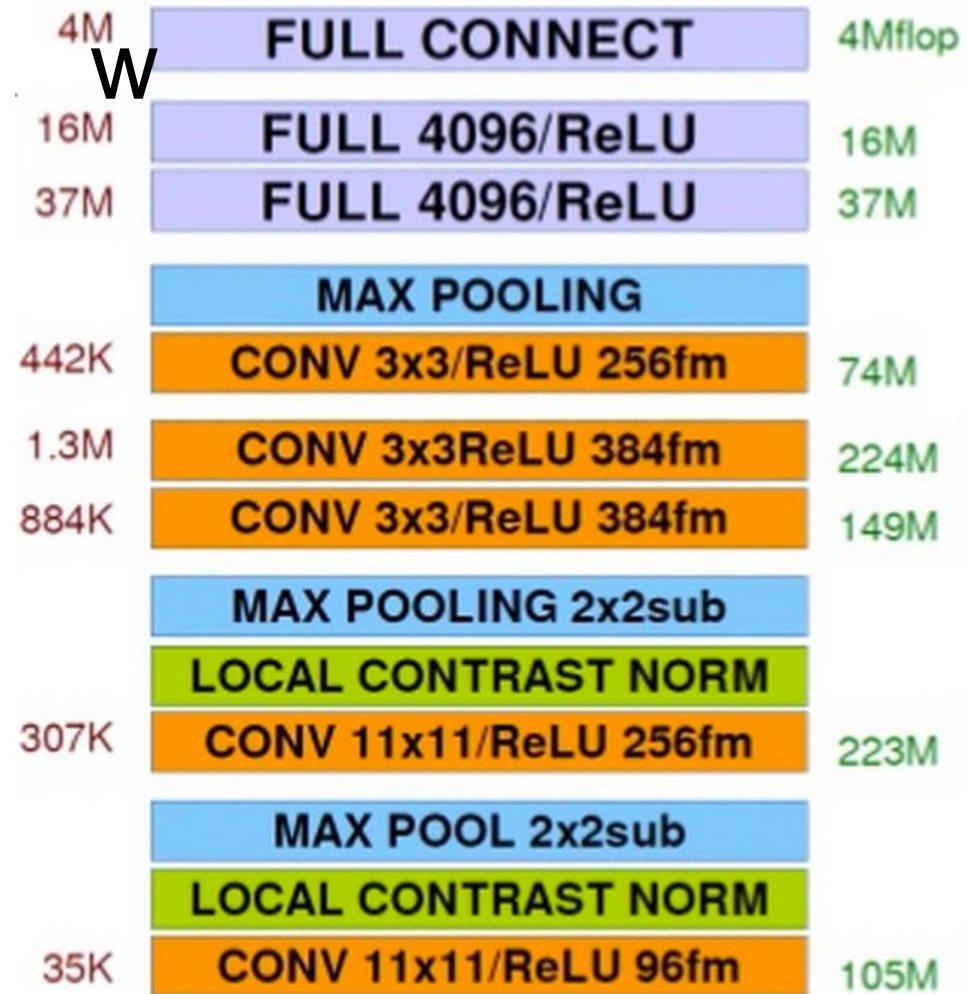


- $s < z \rightarrow$  overlapping pooling
- top-1 and top-5 error rates decrease by 0.4% and 0.3%, respectively, compared to the non-overlapping scheme  $s = 2, z = 2$

# Architecture



# Architecture Overview





# Reducing Overfitting



## Data Augmentation

- 60 million parameters, 650,000 neurons  
→ Overfits a lot.
- Crop 224x224 patches (and their horizontal reflections.)

# Reducing Overfitting

## Data Augmentation

- At test time, average the predictions on the 10 patches.



# Reducing

- Softmax

$$L = \frac{1}{N} \sum_i -\log \left( \frac{e^{f_{y_i}}}{\sum_j e^{f_j}} \right) + \lambda \sum_k \sum_l W_{k,l}^2$$

$j = 1 \dots 1000$

$P(y_i | x_i; W)$  Likelihood

- No need to calibrate to average the predictions over 10 patches.

*cf.* SVM

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} \left[ \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + \Delta) + \lambda \sum_k \sum_l W_{k,l}^2 \right]$$

# Reducing Overfitting

## Data Augmentation

- Change the intensity of RGB channels
- 

$$I_{xy} = [I_{xy}, I_{xy}^G, I_{xy}^{BR}]^T$$


add multiples of principle  
components

$$[\mathbf{p}_1, \mathbf{p}_2, \mathbf{p}_3][\alpha_1 \lambda_1, \alpha_2 \lambda_2, \alpha_3 \lambda_3]^T$$


$$\langle_i \sim N(0, 0.1)$$

# Reducing Overfitting

Dropout



Standard Neural Net



After applying dropout.

- With probability 0.5
- last two 4096 fully-connected layers.



# Stochastic Gradient Descent Learning

## Momentum Update

$$\begin{aligned} v_{i+1} &:= \underbrace{0.9}_{\text{momentum(damping parameter)}} \cdot v_i - \underbrace{0.0005}_{\text{weight decay}} \cdot \epsilon \cdot w_i - \underbrace{\epsilon}_{\text{Learning rate (initialized at 0.01)}} \cdot \underbrace{\left\langle \frac{\partial L}{\partial w} \Big|_{w_i} \right\rangle_{D_i}}_{\substack{\text{Gradient of Loss} \\ \text{w.r.t weight} \\ \text{Averaged over batch}}} \\ w_{i+1} &:= w_i + v_{i+1} \end{aligned}$$

Batch size: 128

- The training took **5 to 6 days** on **two NVIDIA GTX 580 3GB GPUs**.

# Results: ILSVRC-2010

Model	Top-1	Top-5
<i>Sparse coding [2]</i>	47.1%	28.2%
<i>SIFT + FVs [24]</i>	45.7%	25.7%
CNN	<b>37.5%</b>	<b>17.0%</b>

Table 1: Comparison of results on ILSVRC-2010 test set. In *italics* are best results achieved by others.

# Results: ILSVRC-2012

Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs [7]</i>	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	<b>16.4%</b>
1 CNN*	39.0%	16.6%	—
7 CNNs*	36.7%	15.4%	<b>15.3%</b>

Table 2: Comparison of error rates on ILSVRC-2012 validation and test sets. In *italics* are best results achieved by others. Models with an asterisk\* were “pre-trained” to classify the entire ImageNet 2011 Fall release. See Section 6 for details.

# 96 Convolutional



- 11 x 11 x 3 size kernels.
- top 48 kernels on GPU 1 : color-agnostic
- bottom 48 kernels on GPU 2 : color-specific.

Why?

# Eight ILSVRC-2010 test images





# Five ILSVRC-2010 test



The output from the last 4096 fully-connected layer  
: 4096 dimensional feature.

# Discussion

- Depth is really important.

removing a single convolutional layer degrades the performance.

*K. Simonyan, A. Zisserman.*

[Very Deep Convolutional Networks for Large-Scale Image Recognition](#). Technical report, 2014.

→ 16-layer model, 19-layer model. 7.3% top-5 test error on ILSVRC-2012